


Working with tidymodels

OCRUG meetup

Emil Hvitfeldt

2019-1-29

tidymodels is a "meta-package" for modeling and statistical analysis that share the underlying design philosophy, grammar, and data structures of the tidyverse.



```
library(tidymodels)
```

```
## ✓ broom      0.5.1      ✓ purrr      0.3.0
## ✓ dials      0.0.2      ✓ recipes    0.1.4
## ✓ dplyr      0.7.8      ✓ rsample    0.0.4
## ✓ ggplot2    3.1.0      ✓ tibble     2.0.1
## ✓ infer      0.4.0      ✓ yardstick  0.0.2
## ✓ parsnip    0.0.1.9000
```

The packages

- broom
- dials
- dplyr
- ggplot2
- infer
- parsnip
- purrr
- recipes
- rsample
- tibble
- yardstick

The packages (tidyverse)

- broom
- dials
- **dplyr**
- **ggplot2**
- infer
- parsnip
- **purrr**
- recipes
- rsample
- **tibble**
- yardstick

The packages (tidyverse)

- broom
- dials
- **dplyr**
- **ggplot2**
- infer
- parsnip
- **purrr**
- recipes
- rsample
- **tibble**
- yardstick

The packages

- broom
- dials
- dplyr
- ggplot2
- infer
- parsnip
- purrr
- recipes
- rsample
- tibble
- yardstick

Disclaimer

This talk is not designed to give opinions with respect to modeling best practices.

This talk is designed to showcase what packages are available and what they can do.

Consider 32 cars from 1973-74

```
head(mtcars)
```

##		mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
##	Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
##	Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
##	Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
##	Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
##	Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
##	Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1


```
model_glm <- glm(am ~ disp + drat + qsec, data = mtcars,  
                 family = "binomial")
```

```
model_glm <- glm(am ~ disp + drat + qsec, data = mtcars,
                 family = "binomial")
```

```
predict(model_glm)
```

##	Mazda RX4	Mazda RX4 Wag	Datsun 710
##	105.97448	69.85214	27.10173
##	Hornet 4 Drive	Hornet Sportabout	Valiant
##	-276.59440	-240.20025	-313.42683
##	Duster 360	Merc 240D	Merc 230
##	-158.99087	-123.81721	-284.08255
##	Merc 280	Merc 280C	Merc 450SE
##	-20.37716	-59.07966	-167.78204
##	Merc 450SL	Merc 450SLC	Cadillac Fleetwood
##	-180.68287	-206.48454	-458.77204
##	Lincoln Continental	Chrysler Imperial	Fiat 128
##	-427.72554	-357.75792	27.25037
##	Honda Civic	Toyota Corolla	Toyota Corona
##	164.39652	20.76278	-90.84576
##	Dodge Challenger	AMC Javelin	Camaro Z28
##	-211.90064	-189.27739	-74.78345
##	Pontiac Firebird	Fiat X1-9	Porsche 914-2
##	-297.35324	63.64819	184.39936
##	Lotus Europa	Ford Pantera L	Ferrari Dino
##	146.50220	24.28831	162.60217
##	Maserati Bora	Volvo 142E	
##	21.69348	33.80864	

```
library(glmnet)
model_glmnet <- glmnet(am ~ disp + drat + qsec, data = mtcars,
                      family = "binomial")
```

```
library(glmnet)
model_glmnet <- glmnet(am ~ disp + drat + qsec, data = mtcars,
                      family = "binomial")
```

```
## Error in glmnet(am ~ disp + drat + qsec, data = mtcars, family = "binomial"): unused argument (data
```

```
library(glmnet)
model_glmnet <- glmnet(am ~ disp + drat + qsec, data = mtcars,
                      family = "binomial")
```

```
## Error in glmnet(am ~ disp + drat + qsec, data = mtcars, family = "binomial"): unused argument (data
```

```
model_glmnet <- glmnet(x = as.matrix(mtcars[, c("disp", "drat", "qsec")]),
                      y = mtcars[, "am"],
                      family = "binomial")
```

```
library(glmnet)
model_glmnet <- glmnet(am ~ disp + drat + qsec, data = mtcars,
                      family = "binomial")
```

```
## Error in glmnet(am ~ disp + drat + qsec, data = mtcars, family = "binomial"): unused argument (data
```

```
model_glmnet <- glmnet(x = as.matrix(mtcars[, c("disp", "drat", "qsec")]),
                      y = mtcars[, "am"],
                      family = "binomial")
```

```
model_glm <- glm(x = as.matrix(mtcars[, c("disp", "drat", "qsec")]),
                y = mtcars[, "am"],
                family = "binomial")
```

```
library(glmnet)
model_glmnet <- glmnet(am ~ disp + drat + qsec, data = mtcars,
                      family = "binomial")
```

```
## Error in glmnet(am ~ disp + drat + qsec, data = mtcars, family = "binomial"): unused argument (data
```

```
model_glmnet <- glmnet(x = as.matrix(mtcars[, c("disp", "drat", "qsec")]),
                      y = mtcars[, "am"],
                      family = "binomial")
```

```
model_glm <- glm(x = as.matrix(mtcars[, c("disp", "drat", "qsec")]),
                y = mtcars[, "am"],
                family = "binomial")
```

```
## Error in environment(formula): argument "formula" is missing, with no default
```

User-facing problems in modeling in R

- Data must be a matrix (except when it needs to be a data.frame)
- Must use formula or x/y (or both)
- Inconsistent naming of arguments (ntree in randomForest, num.trees in ranger)
- na.omit explicitly or silently
- May or may not accept factors

User-facing problems in modeling in R

- Data must be a matrix (except when it needs to be a data.frame)
- Must use formula or x/y (or both)
- Inconsistent naming of arguments (ntree in randomForest, num.trees in ranger)
- na.omit explicitly or silently
- May or may not accept factors



Syntax for Computing Predicted Class Probabilities

Function	Package	Code
lda	MASS	<code>predict(obj)</code>
glm	stats	<code>predict(obj, type = "response")</code>
gbm	gbm	<code>predict(obj, type = "response", n.trees)</code>
mda	mda	<code>predict(obj, type = "posterior")</code>
rpart	rpart	<code>predict(obj, type = "prob")</code>
Weka	RWeka	<code>predict(obj, type = "probability")</code>
logitboost	LogitBoost	<code>predict(obj, type = "raw", nIter)</code>



The goals of **parsnip** is...

- Decouple the **model classification** from the **computational engine**
- Separate the definition of a model from its evaluation
- Harmonize argument names
- Make consistent predictions (always tibbles with na.omit=FALSE)

```
model_glm <- glm(am ~ disp + drat + qsec, data = mtcars,  
                 family = "binomial")
```

```
library(parsnip)
model_glm <- logistic_reg(mode = "classification") %>%
  set_engine("glm")

model_glm
```

```
## Logistic Regression Model Specification (classification)
##
## Computational engine: glm
```

```
library(parsnip)
model_glm <- logistic_reg(mode = "classification") %>%
  set_engine("glm")

model_glm
```

```
## Logistic Regression Model Specification (classification)
##
## Computational engine: glm
```

```
fit_glm <- model_glm %>%
  fit(factor(am) ~ disp + drat + qsec, data = mtcars)
```

```
library(parsnip)
model_glmnet <- logistic_reg(mode = "classification") %>%
  set_engine("glmnet")
model_glmnet
```

```
## Logistic Regression Model Specification (classification)
##
## Computational engine: glmnet
```

```
fit_glmnet <- model_glmnet %>%
  fit(factor(am) ~ disp + drat + qsec, data = mtcars)
```


Using both formula and x/y

Formula

```
fit_glm <- model_glm %>%  
  fit(factor(am) ~ ., data = mtcars)
```

x/y

```
fit_glm <- model_glm %>%  
  fit_xy(x = as.matrix(mtcars[, c("disp", "drat", "qsec")]),  
        y = factor(mtcars[, "am"]),  
        data = mtcars)
```

Tidy prediction

```
predict(fit_glm, mtcars)
```

```
## # A tibble: 32 x 1
##   .pred_class
##   <fct>
## 1 1
## 2 1
## 3 1
## 4 0
## 5 0
## 6 0
## 7 0
## 8 0
## 9 0
## 10 0
## # ... with 22 more rows
```

Consider now that we wanted to model a more advanced relationship between variables

```
fit_glm <- model_glm %>%  
  fit(factor(am) ~ poly(mpg, 3) + pca(displ:wt)[1] + pca(displ:wt)[2] + pca(displ:wt)[3],  
      data = mtcars)
```

Consider now that we wanted to model a more advanced relationship between variables

```
fit_glm <- model_glm %>%  
  fit(factor(am) ~ poly(mpg, 3) + pca(displacement)[1] + pca(displacement)[2] + pca(displacement)[3],  
      data = mtcars)
```

- Not all inline functions can be used with formulas
- Having to run some calculations many many times
- Connected to the model, calculations are not saved between models

Post by Max Kuhn about the bad sides of formula

<https://rviews.rstudio.com/2017/03/01/the-r-formula-method-the-bad-parts/>



Preprocessing steps

Some of things you may need to deal with before you can start modeling

- Same unit (center and scale)
- Remove correlation (filter and PCA extraction)
- Missing data (imputation)
- Dummy variables
- Zero Variance

Same units

```
library(recipes)
car_rec <- recipe(mpg ~ ., mtcars) %>%
  step_center(all_predictors()) %>%
  step_scale(all_predictors())
```

PCA

```
library(recipes)
car_rec <- recipe(mpg ~ ., mtcars) %>%
  step_pca(all_predictors(), threshold = 0.8)
```

Any combination of steps

```
car_rec <- recipe(mpg ~ ., mtcars) %>%
  step_knnimpute(drat, wt, neighbors = 5) %>%
  step_zv(all_predictors()) %>%
  step_pca(all_predictors(), threshold = 0.8)
```

```
library(recipes)
car_rec <- recipe(mpg ~ ., mtcars) %>%
  step_center(all_predictors()) %>%
  step_scale(all_predictors())
```



```
library(recipes)
car_rec <- recipe(mpg ~ ., mtcars) %>%
  step_center(all_predictors()) %>%
  step_scale(all_predictors())
```

```
car_rec
```

```
## Data Recipe
##
## Inputs:
##
##      role #variables
## outcome      1
## predictor     10
##
## Operations:
##
## Centering for all_predictors()
## Scaling for all_predictors()
```

```
library(recipes)
car_rec <- recipe(mpg ~ ., mtcars) %>%
  step_center(all_predictors()) %>%
  step_scale(all_predictors())

car_preped <- prep(car_rec, training = mtcars)
```

```
library(recipes)
car_rec <- recipe(mpg ~ ., mtcars) %>%
  step_center(all_predictors()) %>%
  step_scale(all_predictors())

car_preped <- prep(car_rec, training = mtcars)
```

```
bake(car_preped, new_data = mtcars)
```

```
library(recipes)
car_rec <- recipe(mpg ~ ., mtcars) %>%
  step_center(all_predictors()) %>%
  step_scale(all_predictors())

car_prepd <- prep(car_rec, training = mtcars)
```

```
bake(car_prepd, new_data = mtcars)
```

```
## # A tibble: 32 x 11
##   mpg     cyl  disp    hp  drat    wt    qsec    vs    am  gear
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  21    -0.105 -0.571 -0.535  0.568 -0.610 -0.777 -0.868  1.19  0.424
## 2  21    -0.105 -0.571 -0.535  0.568 -0.350 -0.464 -0.868  1.19  0.424
## 3 22.8   -1.22  -0.990 -0.783  0.474 -0.917  0.426  1.12  1.19  0.424
## 4 21.4   -0.105  0.220 -0.535 -0.966 -0.00230 0.890  1.12 -0.814 -0.932
## 5 18.7    1.01   1.04   0.413 -0.835  0.228  -0.464 -0.868 -0.814 -0.932
## 6 18.1   -0.105 -0.0462 -0.608 -1.56   0.248   1.33  1.12 -0.814 -0.932
## 7 14.3    1.01   1.04   1.43  -0.723  0.361  -1.12 -0.868 -0.814 -0.932
## 8 24.4   -1.22  -0.678 -1.24   0.175 -0.0278 1.20  1.12 -0.814  0.424
## 9 22.8   -1.22  -0.726 -0.754  0.605 -0.0687 2.83  1.12 -0.814  0.424
## 10 19.2  -0.105 -0.509 -0.345  0.605  0.228  0.253  1.12 -0.814  0.424
## # ... with 22 more rows, and 1 more variable: carb <dbl>
```

```
library(recipes)
car_rec <- recipe(mpg ~ ., mtcars) %>%
  step_center(all_predictors()) %>%
  step_scale(all_predictors())

car_prepred <- prep(car_rec, training = mtcars)
```

```
juice(car_prepred)
```

```
## # A tibble: 32 x 11
##       cyl    disp    hp  drat    wt    qsec    vs    am    gear  carb
##   <dbl>  <dbl>  <dbl> <dbl>  <dbl> <dbl>  <dbl> <dbl> <dbl> <dbl>
## 1 -0.105 -0.571 -0.535  0.568 -0.610 -0.777 -0.868  1.19  0.424  0.735
## 2 -0.105 -0.571 -0.535  0.568 -0.350 -0.464 -0.868  1.19  0.424  0.735
## 3 -1.22  -0.990 -0.783  0.474 -0.917  0.426  1.12  1.19  0.424 -1.12
## 4 -0.105  0.220 -0.535 -0.966 -0.00230  0.890  1.12 -0.814 -0.932 -1.12
## 5  1.01    1.04   0.413 -0.835  0.228 -0.464 -0.868 -0.814 -0.932 -0.503
## 6 -0.105 -0.0462 -0.608 -1.56  0.248  1.33  1.12 -0.814 -0.932 -1.12
## 7  1.01    1.04   1.43  -0.723  0.361 -1.12 -0.868 -0.814 -0.932  0.735
## 8 -1.22  -0.678 -1.24   0.175 -0.0278  1.20  1.12 -0.814  0.424 -0.503
## 9 -1.22  -0.726 -0.754  0.605 -0.0687  2.83  1.12 -0.814  0.424 -0.503
## 10 -0.105 -0.509 -0.345  0.605  0.228  0.253  1.12 -0.814  0.424  0.735
## # ... with 22 more rows, and 1 more variable: mpg <dbl>
```

```
recipe -> prepare -> bake/juice  
(define) -> (estimate) -> (apply)
```

Types of data splitting

- Random
- By date
- By outcome
 - Classification: within class
 - regression: within quantile

Training and Testing sets

```
library(rsample)

car_preped <- prep(car_rec, training = mtcars)
```




Training and Testing sets

```
library(rsample)

set.seed(4595)

# These slides were almost finished and I didn't want to change the data in all the other slides
big_mtcars <- rerun(10, mtcars) %>%
  bind_rows()

data_split <- initial_split(big_mtcars, strata = "mpg", p = 0.80)

# Training and test data
cars_train <- training(data_split)
cars_test  <- testing(data_split)

car_prep <- prep(car_rec, training = cars_train)

# Preprocessed data
cars_train_p <- juice(car_prep)
cars_test_p  <- bake(car_prep, new_data = cars_test)
```

Cross-Validating (sneak peak)

```
set.seed(1234)
cv_splits <- vfold_cv(
  data = big_mtcars,
  v = 10,
  strata = "mpg"
)
```

```
cv_splits
```

```
## # 10-fold cross-validation using stratification
## # A tibble: 10 x 2
##   splits      id
##   <list>    <chr>
## 1 <split [288/32]> Fold01
## 2 <split [288/32]> Fold02
## 3 <split [288/32]> Fold03
## 4 <split [288/32]> Fold04
## 5 <split [288/32]> Fold05
## 6 <split [288/32]> Fold06
## 7 <split [288/32]> Fold07
## 8 <split [288/32]> Fold08
## 9 <split [288/32]> Fold09
## 10 <split [288/32]> Fold10
```

```

car_form <- mpg ~ disp + qsec + cyl
# Fit on a single analysis resample
fit_model <- function(split, spec) {
  fit(
    object = nearest_neighbor() %>% set_engine("knn"),
    formula = car_form,
    data = analysis(split) # <- pull out training set
  )
}
# For each resample, call fit_model()
cv_splits <- cv_splits %>%
  mutate(models_knn = map(splits, fit_model, spec_lm),
         )
cv_splits

```

```

## # 10-fold cross-validation using stratification
## # A tibble: 10 x 3
##   splits          id      models_knn
##   * <list>        <chr>   <list>
## 1 <split [288/32]> Fold01 <fit[+]>
## 2 <split [288/32]> Fold02 <fit[+]>
## 3 <split [288/32]> Fold03 <fit[+]>
## 4 <split [288/32]> Fold04 <fit[+]>
## 5 <split [288/32]> Fold05 <fit[+]>
## 6 <split [288/32]> Fold06 <fit[+]>
## 7 <split [288/32]> Fold07 <fit[+]>
## 8 <split [288/32]> Fold08 <fit[+]>
## 9 <split [288/32]> Fold09 <fit[+]>
## 10 <split [288/32]> Fold10 <fit[+]>

```



www.rstudio.com

```
library(yardstick)  
head(two_class_example)
```

```
##      truth      Class1      Class2 predicted  
## 1 Class2 0.003589243 0.9964107574      Class2  
## 2 Class1 0.678621054 0.3213789460      Class1  
## 3 Class2 0.110893522 0.8891064779      Class2  
## 4 Class1 0.735161703 0.2648382969      Class1  
## 5 Class2 0.016239960 0.9837600397      Class2  
## 6 Class1 0.999275071 0.0007249286      Class1
```

```
library(yardstick)
head(two_class_example)
```

```
##      truth      Class1      Class2 predicted
## 1 Class2 0.003589243 0.9964107574      Class2
## 2 Class1 0.678621054 0.3213789460      Class1
## 3 Class2 0.110893522 0.8891064779      Class2
## 4 Class1 0.735161703 0.2648382969      Class1
## 5 Class2 0.016239960 0.9837600397      Class2
## 6 Class1 0.999275071 0.0007249286      Class1
```

```
metrics(two_class_example, truth = truth, estimate = predicted)
```

```
## # A tibble: 2 x 3
##   .metric .estimator .estimate
##   <chr>    <chr>         <dbl>
## 1 accuracy binary         0.838
## 2 kap      binary         0.675
```

```
library(yardstick)
head(two_class_example)
```

```
##      truth      Class1      Class2 predicted
## 1 Class2 0.003589243 0.9964107574      Class2
## 2 Class1 0.678621054 0.3213789460      Class1
## 3 Class2 0.110893522 0.8891064779      Class2
## 4 Class1 0.735161703 0.2648382969      Class1
## 5 Class2 0.016239960 0.9837600397      Class2
## 6 Class1 0.999275071 0.0007249286      Class1
```

```
accuracy(two_class_example, truth = truth, estimate = predicted)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>    <chr>         <dbl>
## 1 accuracy binary         0.838
```



```
library(yardstick)
head(two_class_example)
```

```
##      truth      Class1      Class2 predicted
## 1 Class2 0.003589243 0.9964107574      Class2
## 2 Class1 0.678621054 0.3213789460      Class1
## 3 Class2 0.110893522 0.8891064779      Class2
## 4 Class1 0.735161703 0.2648382969      Class1
## 5 Class2 0.016239960 0.9837600397      Class2
## 6 Class1 0.999275071 0.0007249286      Class1
```

```
j_index(two_class_example, truth = truth, estimate = predicted)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 j_index binary      0.673
```

And many more!!

```
library(yardstick)
head(two_class_example)
```

```
##      truth      Class1      Class2 predicted
## 1 Class2 0.003589243 0.9964107574      Class2
## 2 Class1 0.678621054 0.3213789460      Class1
## 3 Class2 0.110893522 0.8891064779      Class2
## 4 Class1 0.735161703 0.2648382969      Class1
## 5 Class2 0.016239960 0.9837600397      Class2
## 6 Class1 0.999275071 0.0007249286      Class1
```

```
conf_mat(two_class_example, truth = truth, estimate = predicted)
```

```
##           Truth
## Prediction Class1 Class2
##      Class1    227     50
##      Class2     31    192
```

```
roc_curve(two_class_example, truth = truth, Class1) %>%  
  autoplot()
```

